



EXTRACT SUMMIT • AUSTIN • 2026

Who's actually buying web data now?

It's the AI companies, and they aren't buying volume.

Andrew Harris

Data acquisition consulting • AB Harris Methods

Former ZoomInfo engineering leader

1B+

pages crawled / month

500M+

SERP queries / month

3,000+

county systems in production

THE QUESTION CHANGED

2019 • WHAT THEY ASKED

“How many records can you get me?”

2026 • WHAT THEY ASK ME

“What did you miss, and how would you know?”

| *Same sources. Same proxies. Completely different definition of done.*

The old buyer paid for volume

Sales intelligence

Contacts, firmographics,
intent

Price & market intel

Ecommerce, travel, retail

Lead generation

Lists, enrichment, territory

Competitive monitoring

Catalogue, pricing, hiring

A missing row was a rounding error.

That assumption is load-bearing in every system we built.

The new buyers all get graded

01	Foundation labs	Pre- and post-training corpora, domain packs
02	Vertical AI products	RAG over a corpus they have to defend
03	Agent companies	Live tool calls: wrong data, wrong action
04	Eval & safety teams	Benchmarks that have to hold still
05	Intelligence vendors	OSINT and threat intel: provenance is the product

When the output gets graded, every gap in the input becomes a liability.

The first call now asks for proof

THE OLD BUYER

- How many records?
- How often do you refresh?
- What's the cost per record?
- Can I get it as a CSV?
- Which fields do you cover?

THE NEW BUYER

- What's your coverage denominator?
- How stale is your oldest row?
- Where did this field come from?
- Will the corpus move between my evals?
- When a source breaks, who finds out?

The left column you answer from a dashboard.

The right column you answer only from architecture.

PART 1

The four things they actually ask for

1 · Coverage

2 · Freshness

3 · Provenance

4 · Reproducibility

Coverage is a denominator, not a success rate

“You crawl 3,000 counties. Which ones were missing last month, and is the gap random?”

WHY THE OLD STACK MISSES IT

98.7% counts only the fetches we tried. If the 40 counties that dropped are all rural, the model learns a bias no eval will catch.

WHAT TO CHANGE

You already have the roster. Ship it, with the exclusion list, so buyers see the shape of the gap, not just the rate.

Freshness is a distribution, not a schedule

“Your SLA says monthly. What's the p99 age of the rows in the file you just sent me?”

WHY THE OLD STACK MISSES IT

The run finished yesterday. The row didn't: when a fetch fails, the old value rides into the new file looking current.

WHAT TO CHANGE

Report age per source, at the tail, from when each value was last confirmed, not last attempted.

Provenance has to survive the delivery

“This says 400 employees. Show me the page that said so, as it looked when you read it.”

WHY THE OLD STACK MISSES IT

Most teams keep raw HTML. Few can trace a shipped field back to its exact snapshot and parser, months later.

WHAT TO CHANGE

Put that link in the record itself, and treat raw-page retention as a contract term, not a storage setting.

Reproducibility is the one nobody prices in

“If my benchmark score moves next quarter, was that my model or your data?”

WHY IT'S THE HARDEST ONE

A selector change shifts 3% of rows. Counts look normal. The buyer spends two weeks blaming their model.

WHAT TO CHANGE

Treat a parser change like a schema migration: diff it against stored pages, then tell buyers what moved.

THE REFRAME

**When someone evaluates a model
on your corpus, you're part of their
experiment control.**

PART 2

Three ways a healthy-looking pipeline lies to you

None of them crash. That's why they reach production.

Three failures your dashboard misses

01

The 200 that means nothing

HTTP 200, valid page, zero results. Dashboard green for five weeks.

Fix: alert on yield per source, not status

02

The run with no identity

“Why did Tuesday's file have 40% fewer rows?” You can only shrug.

Fix: a run ledger you can query

03

The parser that drifted

Null rate creeps up on one field. Row count fine, so nothing alerts.

Fix: field baselines that block the run

Buyers audit everything above the fetch

Publish	Manifest: coverage, counts, QA verdict	AUDITED
QA gate	Can hold a run, not just log it	AUDITED
Delta	What changed per field, what was removed	AUDITED
Normalise	One schema, lineage on every field	AUDITED
Roster & schedule	A declared universe, so gaps are countable	AUDITED
Fetch & escalate	Proxies, browsers, anti-bot, retries	UNSEEN

Buyers never see your fetch layer. They audit the five above it.

Buyers now pay for proof

THE OLD SHAPE

Priced per record

More data wins the deal

Onboarding thrown in free

Churn when a cheaper vendor appears



THE NEW SHAPE

Priced per corpus you can vouch for

Less data with a manifest wins

Mapping and lineage billed as real work

Churn when a benchmark moves

They'll pay more for less data, if you can tell them what they're holding.

Put these in the delivery this quarter

01	Ship the denominator	Roster, attempted, excluded, every delivery	~1 day
02	Report the age tail	p99 per source, from last confirmed	~2 days
03	Link fields to snapshots	Any shipped value traceable to its page	~3 days
04	Publish a parser changelog	What moved when an extractor changed	~2 days
05	Let the gate hold	A bad run waits, and the buyer is told why	~1 week

None of this is research. It's the difference between quotable and procurable.

**The buyer stopped asking how
much you can get.
They're asking what you can
prove.**

Coverage • Freshness • Provenance • Reproducibility. In that order.

Andrew Harris • AB Harris Methods

andrew@abharrismethods.com

abharrismethods.com

Questions?



WEBSITE



LINKEDIN

FREE FEASIBILITY CHECK

Send me one source you're stuck
on. One-page read back in 48
hours.